

Predicción de Churn usando Random Forest en R

AUTHOR

Nestor Montano P

Caso: Customer Churn Prediction 2020

Esta competencia consiste en predecir si un cliente cambiará de proveedor de telecomunicaciones, algo que se conoce como "Churn".

Kostas Diamantaras. (2020). Customer Churn Prediction 2020. Kaggle.

<https://kaggle.com/competitions/customer-churn-prediction-2020>

EDA y Preliminares

Librerías

```
## R
library(tidyverse) # Conjunto de paquetes para manejo de datos
library(magrittr) # Pipe
library(tidymodels) # Machine Learning en R
library(skimr) # Descriptivas univariadas masivas
library(ranger) # Random Forest
## Estos son para hacer computacion en paralelo en Windows
library(parallel)
library(doParallel)
```

Importar Datos

Importar y modificar los objetos para que sean del tipo correcto

```
# Leer el archivo de excel y asignarlo al objeto data
data <- read_csv(file = "Data/train_kaggle.csv")
```

Rows: 4250 Columns: 20

-- Column specification -----

Delimiter: ","

chr (5): state, area_code, international_plan, voice_mail_plan, churn

dbl (15): account_length, number_vmail_messages, total_day_minutes, total_da...

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

```
data %>% glimpse
```

Rows: 4,250

Columns: 20

\$ state <chr> "OH", "NJ", "OH", "OK", "MA", "MO", "LA"~

\$ account_length <dbl> 107, 137, 84, 75, 121, 147, 117, 141, 65~

```
$ area_code           <chr> "area_code_415", "area_code_415", "area_~
$ international_plan  <chr> "no", "no", "yes", "yes", "no", "yes", "~
$ voice_mail_plan     <chr> "yes", "no", "no", "no", "yes", "no", "n~
$ number_vmail_messages <dbl> 26, 0, 0, 0, 24, 0, 0, 37, 0, 0, 0, 0, 0~
$ total_day_minutes   <dbl> 161.6, 243.4, 299.4, 166.7, 218.2, 157.0~
$ total_day_calls     <dbl> 123, 114, 71, 113, 88, 79, 97, 84, 137, ~
$ total_day_charge    <dbl> 27.47, 41.38, 50.90, 28.34, 37.09, 26.69~
$ total_eve_minutes   <dbl> 195.5, 121.2, 61.9, 148.3, 348.5, 103.1,~
$ total_eve_calls     <dbl> 103, 110, 88, 122, 108, 94, 80, 111, 83,~
$ total_eve_charge    <dbl> 16.62, 10.30, 5.26, 12.61, 29.62, 8.76, ~
$ total_night_minutes <dbl> 254.4, 162.6, 196.9, 186.9, 212.6, 211.8~
$ total_night_calls   <dbl> 103, 104, 89, 121, 118, 96, 90, 97, 111,~
$ total_night_charge  <dbl> 11.45, 7.32, 8.86, 8.41, 9.57, 9.53, 9.7~
$ total_intl_minutes  <dbl> 13.7, 12.2, 6.6, 10.1, 7.5, 7.1, 8.7, 11~
$ total_intl_calls    <dbl> 3, 5, 7, 3, 7, 6, 4, 5, 6, 5, 2, 5, 9, 4~
$ total_intl_charge   <dbl> 3.70, 3.29, 1.78, 2.73, 2.03, 1.92, 2.35~
$ number_customer_service_calls <dbl> 1, 0, 2, 3, 3, 0, 1, 0, 4, 0, 1, 3, 4, 1~
$ churn              <chr> "no", "no", "no", "no", "no", "no", "no"~
```

Corregir tipos de datos

```
# Convertir a factor
data %>%
  mutate( churn = factor(churn,
                        levels= c("yes","no"),
                        labels= c("si", "no"))
  ) -> data
```

EDA Univariado

```
skim(data)
```

Data summary	
Name	data
Number of rows	4250
Number of columns	20
Column type frequency:	
character	4
factor	1
numeric	15
Group variables	
None	

Variable type: character

skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
state	0	1	2	2	0	51	0
area_code	0	1	13	13	0	3	0
international_plan	0	1	2	3	0	2	0
voice_mail_plan	0	1	2	3	0	2	0

Variable type: factor

skim_variable	n_missing	complete_rate	ordered	n_unique	top_counts
churn	0	1	FALSE	2	no: 3652, si: 598

Variable type: numeric

skim_variable	n_missing	complete_rate	mean	sd	p0	p25	p50	p75	p100	hist
account_length	0	1	100.24	39.70	1	73.00	100.00	127.00	243.00	
number_vmail_messages	0	1	7.63	13.44	0	0.00	0.00	16.00	52.00	
total_day_minutes	0	1	180.26	54.01	0	143.33	180.45	216.20	351.50	
total_day_calls	0	1	99.91	19.85	0	87.00	100.00	113.00	165.00	
total_day_charge	0	1	30.64	9.18	0	24.36	30.68	36.75	59.76	
total_eve_minutes	0	1	200.17	50.25	0	165.93	200.70	233.78	359.30	
total_eve_calls	0	1	100.18	19.91	0	87.00	100.00	114.00	170.00	
total_eve_charge	0	1	17.02	4.27	0	14.10	17.06	19.87	30.54	
total_night_minutes	0	1	200.53	50.35	0	167.22	200.45	234.70	395.00	
total_night_calls	0	1	99.84	20.09	0	86.00	100.00	113.00	175.00	
total_night_charge	0	1	9.02	2.27	0	7.52	9.02	10.56	17.77	
total_intl_minutes	0	1	10.26	2.76	0	8.50	10.30	12.00	20.00	
total_intl_calls	0	1	4.43	2.46	0	3.00	4.00	6.00	20.00	
total_intl_charge	0	1	2.77	0.75	0	2.30	2.78	3.24	5.40	
number_customer_service_calls	0	1	1.56	1.31	0	1.00	1.00	2.00	9.00	

Balanceo


```
data %>%
  group_by( churn) %>%
  count( name = 'frec') %>%
  ungroup() %>%
  mutate( Porc= frec/sum(frec))
```

A tibble: 2 x 3
 churn frec Porc
 <fct> <int> <dbl>

```
1 si      598 0.141
2 no     3652 0.859
```

```
data %>%
  group_by( churn) %>%
  count( name = 'frec') %>%
  ungroup() %>%
  mutate( Porc= frec/sum(frec)) %>%
  ggplot( aes(x= churn, y= Porc)) +
  geom_segment( aes(xend= churn, y=0, yend=Porc),
                color= "steelblue", linewidth= 1) +
  geom_point( size=5, color= "steelblue") +
  coord_flip() +
  scale_y_continuous( labels = percent_format()) +
  labs(title= 'Porcentaje de Clientes que Abandonan',
        y= "Porcentaje", x= "Churn") +
  theme_bw()
```

Porcentaje de Clientes que Abandonan



Se puede ver una proporción de 6 a 1 entre el "No" y el "Si"

EDA Multivariado

```
# data %>%
#   select_if( is.numeric) %>%
#   GGally::ggscatmat()
```

Otra opción para ver la correlacion es:

```
# data %>%
#   select_if( is.numeric) %>%
#   cor %>%
#   corrrplot::corrplot(
#     method = "number", type = "lower" )
```

Modelamiento

Train-Test Split

```
set.seed(1234) # Semilla para aleatorios
split <- data %>%
  initial_split(
    prop = 0.8, # Porcentaje al train
    strata = churn # Estratificación del muestreo
  )
```

Con el split creado, podemos obtener nuestro train y test, así:

```
train <- training(split)
dim(train)
```

```
[1] 3399  20
```

```
test <- testing(split)
dim(test)
```

```
[1] 851  20
```

Preprocesamiento

Balancear usando pesos

Aqui vamos a crear una columna para la ponderacion de cada fila, esta ponderación se va a usar en la estimación del modelo y en los pasos de la receta que sean "supervisado" (cuando se usa la variable "y" en el preprocesamiento)

```
train %>%
  mutate(
    ## crear la variable con los pesos
    case_wts = ifelse(churn == "si", 6, 1),
    ## crea el vector de importancia ponderada
    case_wts = importance_weights(case_wts)
  ) ->
train
```

Receta de preprocesamiento

```
receta <- train %>%
  recipe(churn ~ . ) %>% ## Crea la receta
  ## Eliminar variables que no usaremos
  # step_rm() %>%
  ## Crear nuevas variables (insight desde el EDA)
  # step_mutate( account_length_anio= account_length/12 )
  ## Imputar los datos
  # step_impute_mean()
  step_impute_knn( all_predictors() ) %>%
  ## Estandarizacion/Normalizacion de numericas
  step_normalize( all_numeric(), -all_outcomes() ) %>%
  ## Crear una categoría "otros" que agrupe a categorías pequeñas
  step_other(all_nominal(), -all_outcomes() , threshold = 0.07, other = "otros") %>%
  ## Crear una categoría "new" para observaciones con labels "no muestreados"
  step_novel(all_nominal(), -all_outcomes() , new_level = "new") %>%
  ## Crear variables indicadoras para cada categoría
  step_dummy(all_nominal(), -all_outcomes() ) %>% # Dummy
  ## Eliminar automáticamente variables con alta correlacion
  ## para evitar la multicolinealidad  $x_i \sim x_j$ 
  # step_corr(all_numeric(), -all_outcomes(), threshold = 0.9) %>%
  ## Tambien podemos eliminar variables con multicolinealidad "a mano"
  step_rm(total_day_charge, total_eve_charge,
          total_night_charge, total_intl_charge) %>% # Eliminar
  ## Eliminar columnas con varianza cercana a cero
  step_nzv(all_predictors())
```

receta

-- Recipe -----

-- Inputs

Number of variables by role

```
outcome:      1
predictor:    19
case_weights: 1
```

-- Operations

- * K-nearest neighbor imputation for: all_predictors()
- * Centering and scaling for: all_numeric(), -all_outcomes()
- * Collapsing factor levels for: all_nominal(), -all_outcomes()
- * Novel factor level assignment for: all_nominal(), -all_outcomes()

- * Dummy variables from: `all_nominal()`, `-all_outcomes()`
- * Variables removed: `total_day_charge`, `total_eve_charge`, ...
- * Sparse, unbalanced variable filter on: `all_predictors()`

Entrenamiento y ajuste de Hiperparámetros

Para ajustar los hiperparámetros usaremos la estrategia de definir mallas de búsqueda y evaluar dichas combinaciones sobre un remuestreo (crossvalidation), para ello necesitamos:

- Remuestreo
- Métricas para evaluar y comparar los modelos

Remuestreo

Se define una estrategia de remuestreo (para poder ajustar hiperparámetros)

```
set.seed(1234)
cv <- vfold_cv(train, v = 5, repeats = 1, strata = churn)
cv
```

```
# 5-fold cross-validation using stratification
# A tibble: 5 x 2
  splits      id
  <list>    <chr>
1 <split [2718/681]> Fold1
2 <split [2719/680]> Fold2
3 <split [2719/680]> Fold3
4 <split [2720/679]> Fold4
5 <split [2720/679]> Fold5
```

Métricas

Así también definimos las métricas que queremos que se ejecuten en cada remuestreo

```
metricas <- metric_set(accuracy, sens, spec, bal_accuracy)
metricas
```

```
# A tibble: 4 x 3
  metric      class      direction
  <chr>      <chr>      <chr>
1 accuracy  class_metric maximize
2 sens      class_metric maximize
3 spec      class_metric maximize
4 bal_accuracy class_metric maximize
```

Especificacion del modelo

Fuente:

* <https://www.tidymodels.org/find/parsnip/>

* https://parsnip.tidymodels.org/reference/rand_forest.html

* https://parsnip.tidymodels.org/reference/details_rand_forest_ranger.html

```
rf_sp <-  
  rand_forest(  
    mtry = tune(), trees = tune(), min_n = tune() ) %>%  
    set_engine("ranger", importance = "impurity") %>%  
    set_mode("classification")
```

Notar que no se ha usado todos los hiperparámetros

workflow

```
rf_wflow <-  
  workflow() %>%  
  add_recipe(receta) %>%  
  add_model(rf_sp) %>%  
  add_case_weights(case_wts) ## Aquí agregamos los pesos  
  
rf_wflow
```

== Workflow =====

Preprocessor: Recipe

Model: rand_forest()

-- Preprocessor -----

7 Recipe Steps

- * step_impute_knn()
- * step_normalize()
- * step_other()
- * step_novel()
- * step_dummy()
- * step_rm()
- * step_nzv()

-- Case Weights -----

case_wts

-- Model -----

Random Forest Model Specification (classification)

Main Arguments:

- mtry = tune()
- trees = tune()
- min_n = tune()

Engine-Specific Arguments:

- importance = impurity

Computational engine: ranger

Afinamiento de hiperparametros

Malla de Busqueda

```
set.seed(123)
rf_grid <- rf_sp %>%
  ## preguntamos los parametros tuneables del modelo
  parameters() %>%
  ## Vamos a definir un rango para el min_n y mtry
  update(min_n= min_n( range= c(70, 170)),
         mtry= mtry( range= c(4, 7))) %>%
  grid_latin_hypercube(size = 10)
```

Warning: `parameters.model_spec()` was deprecated in tune 0.1.6.9003.
i Please use `hardhat::extract_parameter_set_dials()` instead.

Paralelizacion

```
parallel::detectCores(logical=FALSE)
```

```
[1] 4
```

```
cl <- makePSOCKcluster(4)
registerDoParallel(cl)
# parallel::stopCluster(cl) ## Esto se debe ejecutar al final
```

Entrenamiento de Malla de Busqueda en la Crossvalidation

```
set.seed(123)
rf_tuned <- tune_grid(
  rf_wflow, ## Modelo
  resamples= cv, ## Crossvalidation
  grid = rf_grid, ## Malla de Busqueda
  metrics = metricas, ## Metricas
  control= control_grid(allow_par = T, save_pred = T) ## Paralel y Pred
)
rf_tuned
```

```
# Tuning results
# 5-fold cross-validation using stratification
# A tibble: 5 x 5
  splits          id    .metrics          .notes          .predictions
  <list>         <chr> <list>          <list>          <list>
1 <split [2718/681]> Fold1 <tibble [40 x 7]> <tibble [0 x 3]> <tibble>
2 <split [2719/680]> Fold2 <tibble [40 x 7]> <tibble [0 x 3]> <tibble>
3 <split [2719/680]> Fold3 <tibble [40 x 7]> <tibble [0 x 3]> <tibble>
4 <split [2720/679]> Fold4 <tibble [40 x 7]> <tibble [0 x 3]> <tibble>
5 <split [2720/679]> Fold5 <tibble [40 x 7]> <tibble [0 x 3]> <tibble>
```

Evaluemos que tal es cada combinacion segun las principales metricas

```
show_best(rf_tuned, metric = 'accuracy', n = 10)
```

```
# A tibble: 10 x 9
```

	mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	7	773	80	accuracy	binary	0.932	5	0.00469	Preprocessor1_Model~
2	6	467	91	accuracy	binary	0.925	5	0.00645	Preprocessor1_Model~
3	5	325	83	accuracy	binary	0.923	5	0.00719	Preprocessor1_Model~
4	5	1441	106	accuracy	binary	0.919	5	0.00644	Preprocessor1_Model~
5	6	950	130	accuracy	binary	0.914	5	0.00539	Preprocessor1_Model~
6	4	1370	116	accuracy	binary	0.911	5	0.00539	Preprocessor1_Model~
7	5	1904	126	accuracy	binary	0.910	5	0.00557	Preprocessor1_Model~
8	6	1737	158	accuracy	binary	0.909	5	0.00550	Preprocessor1_Model~
9	6	1124	162	accuracy	binary	0.907	5	0.00585	Preprocessor1_Model~
10	5	70	144	accuracy	binary	0.906	5	0.00572	Preprocessor1_Model~

```
show_best(rf_tuned, metric = 'sens', n = 10)
```

```
# A tibble: 10 x 9
```

	mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	6	1737	158	sens	binary	0.856	5	0.0135	Preprocessor1_Model~
2	6	1124	162	sens	binary	0.854	5	0.0137	Preprocessor1_Model~
3	5	1904	126	sens	binary	0.851	5	0.0147	Preprocessor1_Model~
4	5	1441	106	sens	binary	0.851	5	0.0147	Preprocessor1_Model~
5	5	70	144	sens	binary	0.849	5	0.0128	Preprocessor1_Model~
6	4	1370	116	sens	binary	0.849	5	0.0151	Preprocessor1_Model~
7	6	950	130	sens	binary	0.849	5	0.0159	Preprocessor1_Model~
8	6	467	91	sens	binary	0.847	5	0.0109	Preprocessor1_Model~
9	5	325	83	sens	binary	0.847	5	0.0158	Preprocessor1_Model~
10	7	773	80	sens	binary	0.845	5	0.0143	Preprocessor1_Model~

```
show_best(rf_tuned, metric = 'spec', n = 10)
```

```
# A tibble: 10 x 9
```

	mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
	<int>	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	7	773	80	spec	binary	0.946	5	0.00546	Preprocessor1_Model~
2	6	467	91	spec	binary	0.938	5	0.00717	Preprocessor1_Model~
3	5	325	83	spec	binary	0.936	5	0.00737	Preprocessor1_Model~
4	5	1441	106	spec	binary	0.929	5	0.00704	Preprocessor1_Model~
5	6	950	130	spec	binary	0.924	5	0.00656	Preprocessor1_Model~
6	4	1370	116	spec	binary	0.921	5	0.00599	Preprocessor1_Model~
7	5	1904	126	spec	binary	0.920	5	0.00660	Preprocessor1_Model~
8	6	1737	158	spec	binary	0.917	5	0.00643	Preprocessor1_Model~
9	6	1124	162	spec	binary	0.916	5	0.00691	Preprocessor1_Model~
10	5	70	144	spec	binary	0.916	5	0.00768	Preprocessor1_Model~

```
show_best(rf_tuned, metric = 'bal_accuracy', n = 10)
```

```
# A tibble: 10 x 9
```

mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
------	-------	-------	---------	------------	------	---	---------	---------

	<int>	<int>	<int>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	7	773	80	bal_accuracy	binary	0.896	5	0.00707	Preprocessor1_~
2	6	467	91	bal_accuracy	binary	0.893	5	0.00669	Preprocessor1_~
3	5	325	83	bal_accuracy	binary	0.891	5	0.00944	Preprocessor1_~
4	5	1441	106	bal_accuracy	binary	0.890	5	0.00819	Preprocessor1_~
5	6	950	130	bal_accuracy	binary	0.887	5	0.00759	Preprocessor1_~
6	6	1737	158	bal_accuracy	binary	0.886	5	0.00691	Preprocessor1_~
7	5	1904	126	bal_accuracy	binary	0.886	5	0.00729	Preprocessor1_~
8	4	1370	116	bal_accuracy	binary	0.885	5	0.00784	Preprocessor1_~
9	6	1124	162	bal_accuracy	binary	0.885	5	0.00697	Preprocessor1_~
10	5	70	144	bal_accuracy	binary	0.883	5	0.00507	Preprocessor1_~

Malla de Búsqueda

```
set.seed(123)
rf_grid_2 <- crossing(
  min_n = seq(80, 92, 3),
  mtry = c(5, 6),
  trees= seq(500, 800, 100)
)
rf_grid_2
```

```
# A tibble: 40 x 3
  min_n mtry trees
  <dbl> <dbl> <dbl>
1     80     5  500
2     80     5  600
3     80     5  700
4     80     5  800
5     80     6  500
6     80     6  600
7     80     6  700
8     80     6  800
9     83     5  500
10    83     5  600
# i 30 more rows
```

Entrenamiento de Malla de Búsqueda en la Crossvalidation

```
set.seed(123)
rf_tuned_2 <- tune_grid(
  rf_wflow, ## Modelo
  resamples= cv, ## Crossvalidation
  grid = rf_grid_2, ## Malla de Búsqueda
  metrics = metricas, ## Metricas
  control= control_grid(allow_par = T, save_pred = T) ## Paralel y Pred
)
rf_tuned_2
```

```
# Tuning results
# 5-fold cross-validation using stratification
# A tibble: 5 x 5
```

	splits	id	.metrics	.notes	.predictions
	<list>	<chr>	<list>	<list>	<list>
1	<split [2718/681]>	Fold1	<tibble [160 x 7]>	<tibble [0 x 3]>	<tibble>
2	<split [2719/680]>	Fold2	<tibble [160 x 7]>	<tibble [0 x 3]>	<tibble>
3	<split [2719/680]>	Fold3	<tibble [160 x 7]>	<tibble [0 x 3]>	<tibble>
4	<split [2720/679]>	Fold4	<tibble [160 x 7]>	<tibble [0 x 3]>	<tibble>
5	<split [2720/679]>	Fold5	<tibble [160 x 7]>	<tibble [0 x 3]>	<tibble>

Evaluemos que tal es cada combinacion segun las principales metricas

```
show_best(rf_tuned_2, metric = 'accuracy', n = 10)
```

A tibble: 10 x 9

	mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	6	500	83	accuracy	binary	0.930	5	0.00588	Preprocessor1_Model~
2	6	500	80	accuracy	binary	0.930	5	0.00576	Preprocessor1_Model~
3	6	700	80	accuracy	binary	0.929	5	0.00522	Preprocessor1_Model~
4	6	800	80	accuracy	binary	0.929	5	0.00676	Preprocessor1_Model~
5	6	700	83	accuracy	binary	0.929	5	0.00650	Preprocessor1_Model~
6	6	600	80	accuracy	binary	0.929	5	0.00615	Preprocessor1_Model~
7	6	800	83	accuracy	binary	0.929	5	0.00530	Preprocessor1_Model~
8	5	700	80	accuracy	binary	0.928	5	0.00608	Preprocessor1_Model~
9	6	600	83	accuracy	binary	0.927	5	0.00582	Preprocessor1_Model~
10	6	500	89	accuracy	binary	0.927	5	0.00737	Preprocessor1_Model~

```
show_best(rf_tuned_2, metric = 'sens', n = 10)
```

A tibble: 10 x 9

	mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	5	600	80	sens	binary	0.854	5	0.0137	Preprocessor1_Model~
2	5	800	83	sens	binary	0.854	5	0.0137	Preprocessor1_Model~
3	5	800	92	sens	binary	0.854	5	0.0137	Preprocessor1_Model~
4	5	500	83	sens	binary	0.851	5	0.0142	Preprocessor1_Model~
5	5	600	83	sens	binary	0.851	5	0.0117	Preprocessor1_Model~
6	5	500	89	sens	binary	0.851	5	0.0117	Preprocessor1_Model~
7	5	800	89	sens	binary	0.851	5	0.0142	Preprocessor1_Model~
8	5	700	80	sens	binary	0.851	5	0.0147	Preprocessor1_Model~
9	6	500	83	sens	binary	0.851	5	0.0147	Preprocessor1_Model~
10	5	500	86	sens	binary	0.851	5	0.0147	Preprocessor1_Model~

```
show_best(rf_tuned_2, metric = 'spec', n = 10)
```

A tibble: 10 x 9

	mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	6	500	80	spec	binary	0.943	5	0.00662	Preprocessor1_Model~
2	6	500	83	spec	binary	0.943	5	0.00657	Preprocessor1_Model~
3	6	800	80	spec	binary	0.943	5	0.00720	Preprocessor1_Model~
4	6	700	80	spec	binary	0.942	5	0.00569	Preprocessor1_Model~
5	6	600	80	spec	binary	0.942	5	0.00687	Preprocessor1_Model~
6	6	700	83	spec	binary	0.941	5	0.00702	Preprocessor1_Model~

7	6	800	83 spec	binary	0.941	5	0.00566	Preprocessor1_Model~
8	6	600	83 spec	binary	0.941	5	0.00633	Preprocessor1_Model~
9	5	700	80 spec	binary	0.940	5	0.00665	Preprocessor1_Model~
10	6	700	86 spec	binary	0.940	5	0.00674	Preprocessor1_Model~

```
show_best(rf_tuned_2, metric = 'bal_accuracy', n = 10)
```

A tibble: 10 x 9

	mtry	trees	min_n	.metric	.estimator	mean	n	std_err	.config
	<dbl>	<dbl>	<dbl>	<chr>	<chr>	<dbl>	<int>	<dbl>	<chr>
1	6	500	83	bal_accuracy	binary	0.897	5	0.00781	Preprocessor1_~
2	5	600	80	bal_accuracy	binary	0.896	5	0.00814	Preprocessor1_~
3	5	700	80	bal_accuracy	binary	0.896	5	0.00803	Preprocessor1_~
4	6	700	80	bal_accuracy	binary	0.896	5	0.00698	Preprocessor1_~
5	6	500	89	bal_accuracy	binary	0.895	5	0.00850	Preprocessor1_~
6	6	700	83	bal_accuracy	binary	0.895	5	0.00846	Preprocessor1_~
7	6	800	83	bal_accuracy	binary	0.895	5	0.00834	Preprocessor1_~
8	6	800	92	bal_accuracy	binary	0.895	5	0.00810	Preprocessor1_~
9	6	500	80	bal_accuracy	binary	0.895	5	0.00732	Preprocessor1_~
10	5	800	83	bal_accuracy	binary	0.895	5	0.00851	Preprocessor1_~

Bien, podríamos probar mallas más extensas o tomar una decisión ya con las pruebas realizadas.

Modelo final

```
## Definir la mejor combinacion
rf_pars_fin <- select_best(rf_tuned_2, metric = 'sens')

## Finalizar (darle valores a parametros tuneables) el workflow
rf_wflow_fin <-
  rf_wflow %>%
    finalize_workflow(rf_pars_fin)
rf_wflow_fin
```

== Workflow =====

Preprocessor: Recipe

Model: rand_forest()

-- Preprocessor -----

7 Recipe Steps

```
* step_impute_knn()
* step_normalize()
* step_other()
* step_novel()
* step_dummy()
* step_rm()
* step_nzv()
```

-- Case Weights -----

case_wts

-- Model -----

Random Forest Model Specification (classification)

Main Arguments:

```
mtry = 5  
trees = 600  
min_n = 80
```

Engine-Specific Arguments:

```
importance = impurity
```

Computational engine: ranger

Ahora sí, se entrena el modelo final

```
## Entrenar el modelo final  
rf_fitted <- fit(rf_wflow_fin, train)  
rf_fitted
```

== Workflow [trained] =====

Preprocessor: Recipe

Model: rand_forest()

-- Preprocessor -----

7 Recipe Steps

```
* step_impute_knn()  
* step_normalize()  
* step_other()  
* step_novel()  
* step_dummy()  
* step_rm()  
* step_nzv()
```

-- Case Weights -----

case_wts

-- Model -----

Ranger result

Call:

```
ranger::ranger(x = maybe_data_frame(x), y = y, mtry = min_cols(~5, x), num.trees = ~600,  
min.node.size = min_rows(~80, x), importance = ~"impurity", num.threads = 1, verbose =  
FALSE, seed = sample.int(10^5, 1), probability = TRUE, case.weights = weights)
```

Type:	Probability estimation
Number of trees:	600
Sample size:	3399
Number of independent variables:	14
Mtry:	5
Target node size:	80
Variable importance mode:	impurity
Splitrule:	gini
OOB prediction error (Brier s.):	0.06769222

Notar que `arbol_fitted` sigue siendo un workflow, si por algún motivo queremos sólo trabajar con el modelo, podemos:

```
rf_model_fin <- extract_fit_parsnip(rf_fitted)
```

Evaluación del modelo

Vamos a comparar las métricas del modelo en el train como en el test

```
train %>%
  predict(rf_fitted, new_data = .) %>%
  mutate(Real= train$churn) %>%
  conf_mat(truth = Real, estimate = .pred_class) %>%
  summary
```

A tibble: 13 x 3

	.metric	.estimator	.estimate
	<chr>	<chr>	<dbl>
1	accuracy	binary	0.949
2	kap	binary	0.800
3	sens	binary	0.879
4	spec	binary	0.961
5	ppv	binary	0.787
6	npv	binary	0.980
7	mcc	binary	0.802
8	j_index	binary	0.840
9	bal_accuracy	binary	0.920
10	detection_prevalence	binary	0.157
11	precision	binary	0.787
12	recall	binary	0.879
13	f_meas	binary	0.830

```
test %>%
  predict(rf_fitted, new_data = .) %>%
  mutate(Real= test$churn) %>%
  conf_mat(truth = Real, estimate = .pred_class) %>%
  summary
```

A tibble: 13 x 3

	.metric	.estimator	.estimate
	<chr>	<chr>	<dbl>
1	accuracy	binary	0.931
2	kap	binary	0.735
3	sens	binary	0.85
4	spec	binary	0.944
5	ppv	binary	0.713
6	npv	binary	0.975
7	mcc	binary	0.739
8	j_index	binary	0.794
9	bal_accuracy	binary	0.897
10	detection_prevalence	binary	0.168

11 precision	binary	0.713
12 recall	binary	0.85
13 f_meas	binary	0.776

Podemos ver que no existen mucha diferencia, por lo que se puede concluir que el modelo **no se ha sobreajustado**

Finalizar Paralelizacion

```
# parallel::detectCores(logical=FALSE)
# cl <- makePSOCKcluster(4)
# registerDoParallel(cl)
parallel::stopCluster(cl) ## Esto se debe ejecutar al final
```

Análisis Posteriores

¿Qué variables parecen estar más relacionadas con el abandono del cliente?

```
library(vip)
```

Warning: package 'vip' was built under R version 4.0.5

Attaching package: 'vip'

The following object is masked from 'package:utils':

vi

```
rf_model_fin %>%
  vip(geom = "point")
```